



Hardware for Offline AI Agent



All-in-One Edge Agent Dev-kit

Extremely low learning curve - Voice interaction and developer-friendly design, with Model already built-in

Focused on blue-collar scenarios - Tackles challenges arising from limited education, small memory capacity, and restricted work hours

Local integration of search and reasoning - No internet connection required, minimal latency, no data leakage, and ideal for controlled environments

Traction (Since Aug)

3B model edge inference - 7 tokens/second, 5W

Over \$20,000 in revenue

Product buzz surpassing 1 million, with 200+ user interviews conducted

Official Badge Provider for DEFCON 2025

Initiated collaborations with vertical enterprises

Intelligence Explosion

**Agents everywhere
will require a significant amount of hardware carriers**

**Current solutions edge-side Agents are extremely limited
SMB applications are diverse and have high demand
Core data for individuals/enterprises raises privacy concerns**

Intelligence Density

= IQ / Physical Volume

**Today, our box accomplishes tasks equivalent to an IQ of 100,
replacing 0.5 units of labor.**

**In 12 months, next intelligence will exceed 150+, achieving
exponential growth.**

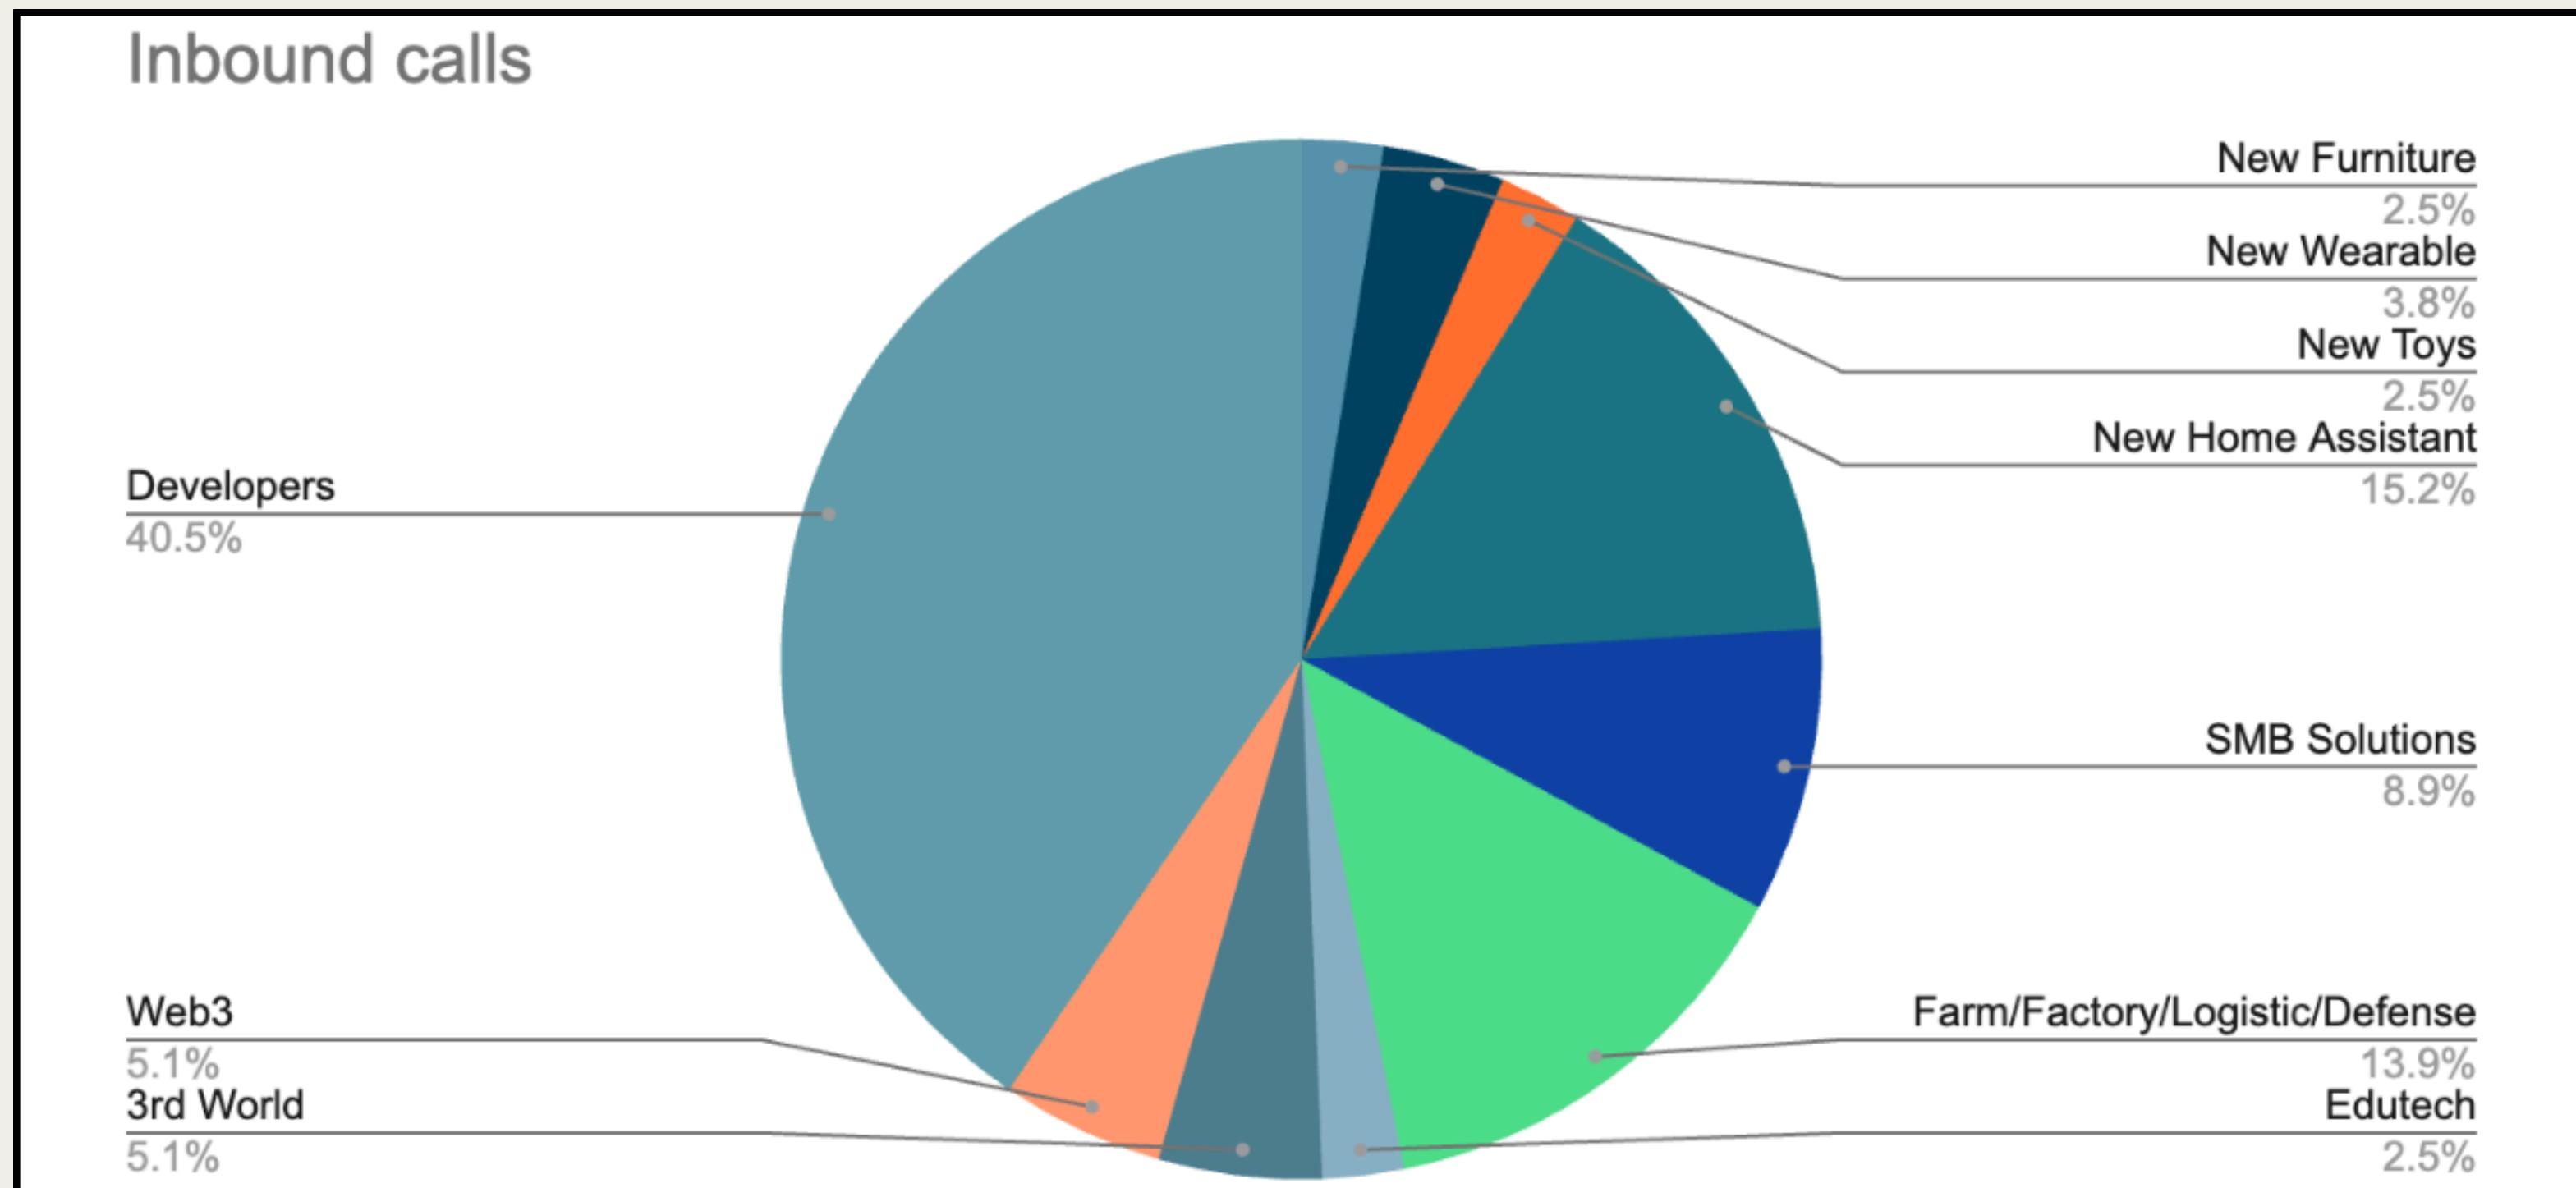
High Demand

November 2024

Sun	Mon	Tue	Wed	Thu	Fri	Sa
27	28	29 • Martin's... 6 AM GMT+8	30 • Foun... 1:45 AM GMT+8 • HardTech Afte... 8 AM	31 Halloween Diwali Halloween 1 more...	Nov 1 • Portfolio... 5 AM GMT+8	2 Day of the Dead
3 Daylight Saving Ti... Daylight Saving Ti... Daylight Saving Ti...	4 • int pamir 1 AM GMT+8	5 Election Day (Gene... Election Day Election Day (Gene...	6 • Tianq... 1:30 AM GMT+8	7 • Portfolio... 6 AM GMT+8	8 • Tush... 3:30 AM GMT+8 • Jason S... 4 AM GMT+8 • Joseph... 8 AM GMT+8	9
10 Rishabh... 3 AM GMT+8 Dre and... 4 AM GMT+8 Jun Ha... 12 PM GMT+8 1 more...	11 Veterans Day Veterans Day Veterans Day 2 more...	12 • Jason S... 2 AM GMT+8 • Sean... 2:30 AM GMT+8 • Patrick... 3 AM GMT+8 8 more...	13 • Leonard... 2 AM GMT+8 • Moe... 2:30 AM GMT+8 • Santhos... 6 AM GMT+8 3 more...	14 • ZHAOQI... 2 AM GMT+8 • 吕旭... 2:30 AM GMT+8 • Alec Ort... 3 AM GMT+8 5 more...	15 • Ryan... 3:30 AM GMT+8 • Nicke... 4:30 AM GMT+8 • Portfolio... 6 AM GMT+8 2 more...	16 • Joseph... 2 AM GMT+8 • Ryan... 2:30 AM GMT+8 • Part... 11:30 AM GMT+8 2 more...
17 Artur an... 3 AM GMT+8 Thas... 3:30 AM GMT+8 Andr... 4:30 AM GMT+8 2 more...	18 • int pamir 1 AM GMT+8 • FEM... 10:30 AM GMT+8 • Dipu... 1:30 PM GMT+8 1 more...	19 • Jonas a... 2 AM GMT+8 • FERN... 2:30 AM GMT+8 • edison... 10 AM GMT+8 3 more...	20 • Ahmed... 4 AM GMT+8 • Tristan... 10 AM GMT+8 • Davi... 10:30 AM GMT+8 4 more...	21 • Riccard... 2 AM GMT+8 • Will... 2:30 AM GMT+8 • NAEEM... 3 AM GMT+8 8 more...	22 • Piyush J... 2 AM GMT+8 • Ivan an... 4 AM GMT+8 • Geor... 4:30 AM GMT+8 2 more...	23 • Daniel R... 2 AM GMT+8 • Alexa... 2:30 AM GMT+8 • Santosh... 3 AM GMT+8 3 more...
24	25 meeting with Jose... • int pamir 1 AM GMT+8 • xuwen oh 3 PM GMT+8	26 Stay at Beijing Marriott Hotel Northeast • Giorg... 2:30 AM GMT+8 • Flight... 2:55 AM GMT+8 3 more...	27 • Tianq... 1:15 AM GMT+8 • Charl... 2:30 AM GMT+8 4 more...	28 Thanksgiving Day • Abdul... 2:30 AM GMT+8 • Thanksgiving 3 more...	29 • Carlos... 2 AM GMT+8 • Abdul... 2:30 AM GMT+8 • Portfolio... 6 AM GMT+8 2 more...	30 • Mitch a... 6 AM GMT+8 • Marq a... 11 AM GMT+8
Dec 1	2 • int pamir 1 AM GMT+8	3 • Patrick... 2 AM GMT+8 • Prab... 2:30 AM GMT+8	4 • Pamir... 2:45 AM GMT+8 • Abdula... 10 AM GMT+8 3 more...	5 • George... 2 AM GMT+8 • Prab... 2:30 AM GMT+8 • Kemar... 3 AM GMT+8 3 more...	6 • Tonee N... 2 AM GMT+8 • Abdul... 2:30 AM GMT+8 • hlynu... 4:30 AM GMT+8 3 more...	7 • Sarah e... 2 AM GMT+8 • Jord... 2:30 AM GMT+8 • Alan an... 3 AM GMT+8 3 more...

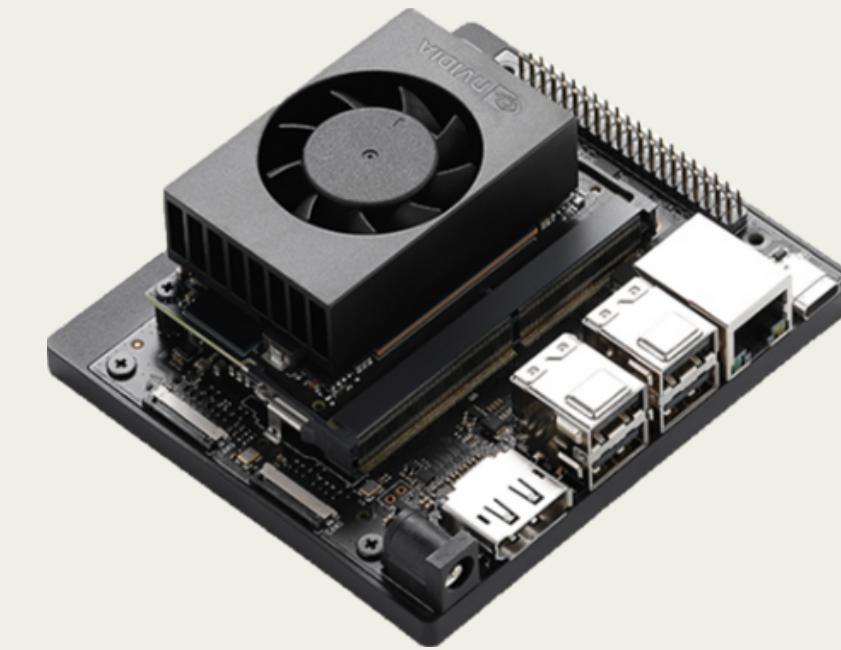
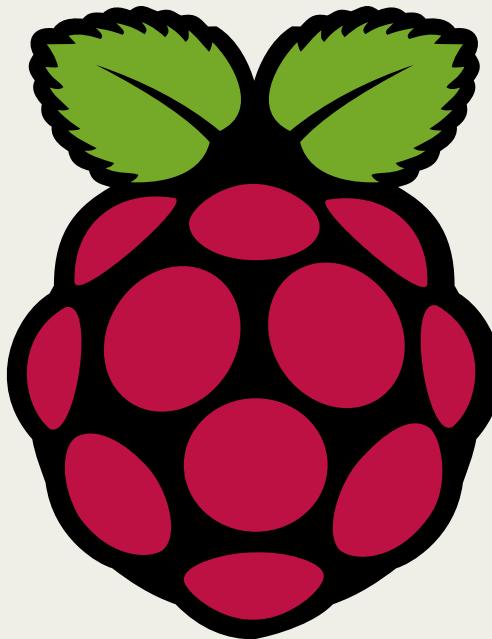
Traction

First batch of product made \$20000+, next iteration demo got us 140+ inbound calls



Real World Use Case

The gap in development platforms for edge Agents

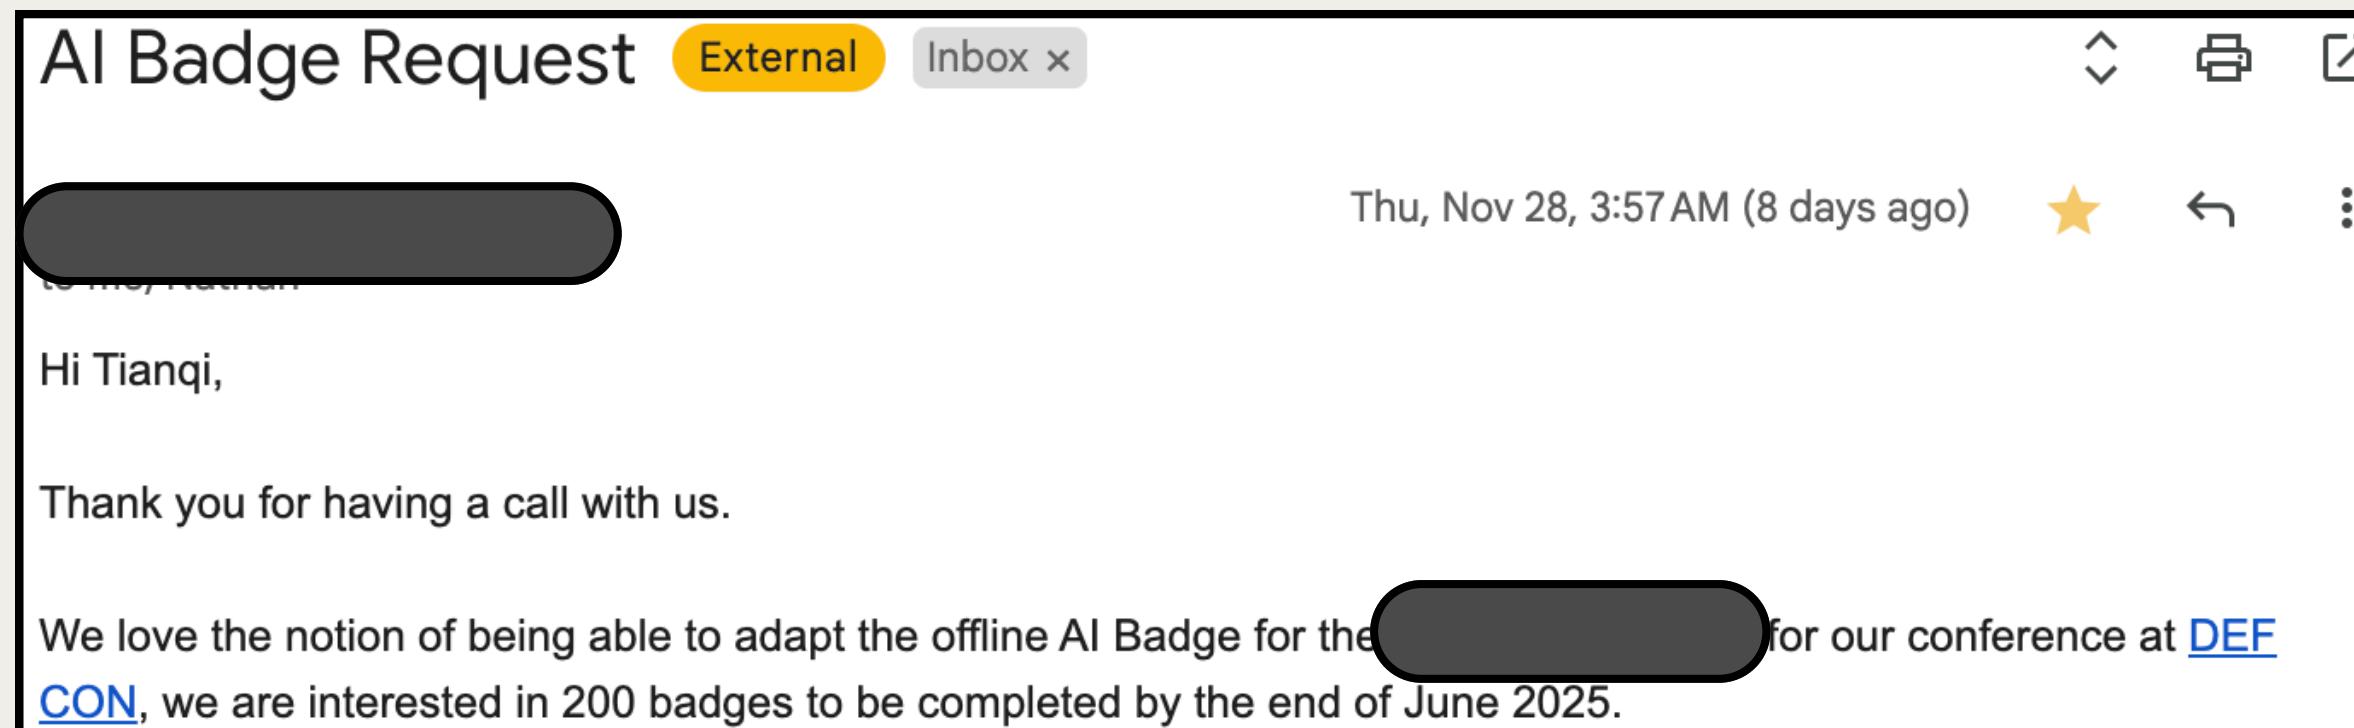


**bring your own solution, where
does edge-side intelligence fit?**

Choice of development platform

The logic is simple: since you have to buy compute power anyway, why not choose one that comes with intelligence pre-packaged?

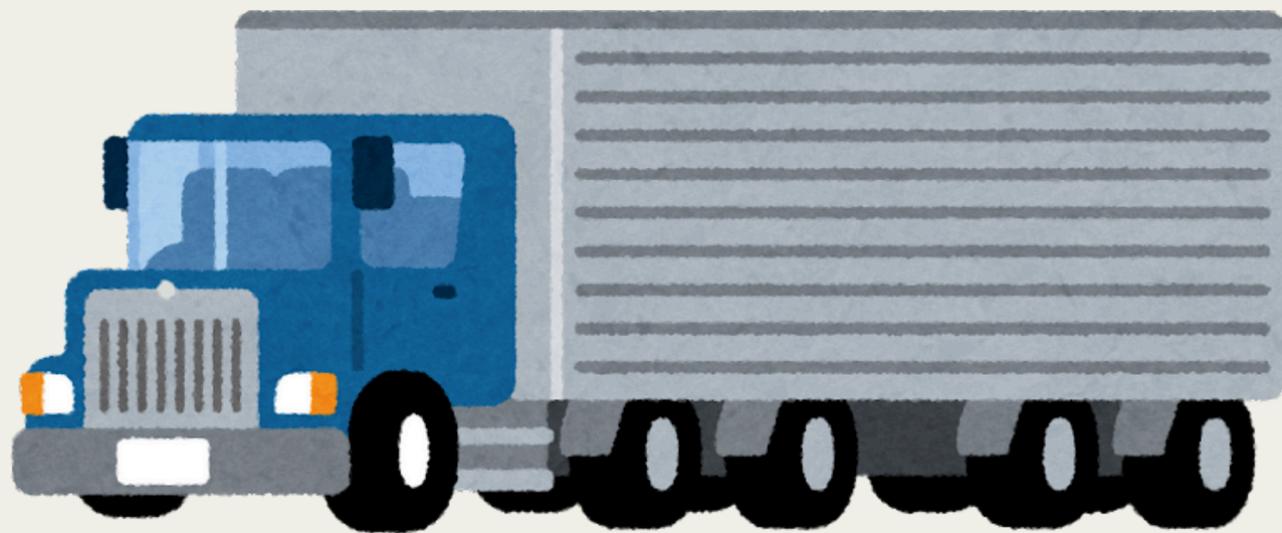
I am from [REDACTED]. We build tiny Gen AI models and edge inference acceleration solutions. Looking for partnership opportunities.



Counterintuitive insight

Selling **intelligence** directly is
better than selling compute power.

Scenario: Large-scale information retrieval in logistics



The top big players in the U.S. logistics industry account for **only 14.93% of the market**.

The remaining market is divided among SMBs. These SMBs lack the capability to develop specialized smart hardware or mobile/desktop apps.

There is also competition among SMBs, and they are actively seeking intelligent alternatives.

Scenario: Large-scale information retrieval in logistics



Is Product A compliant? What information do I need to verify?

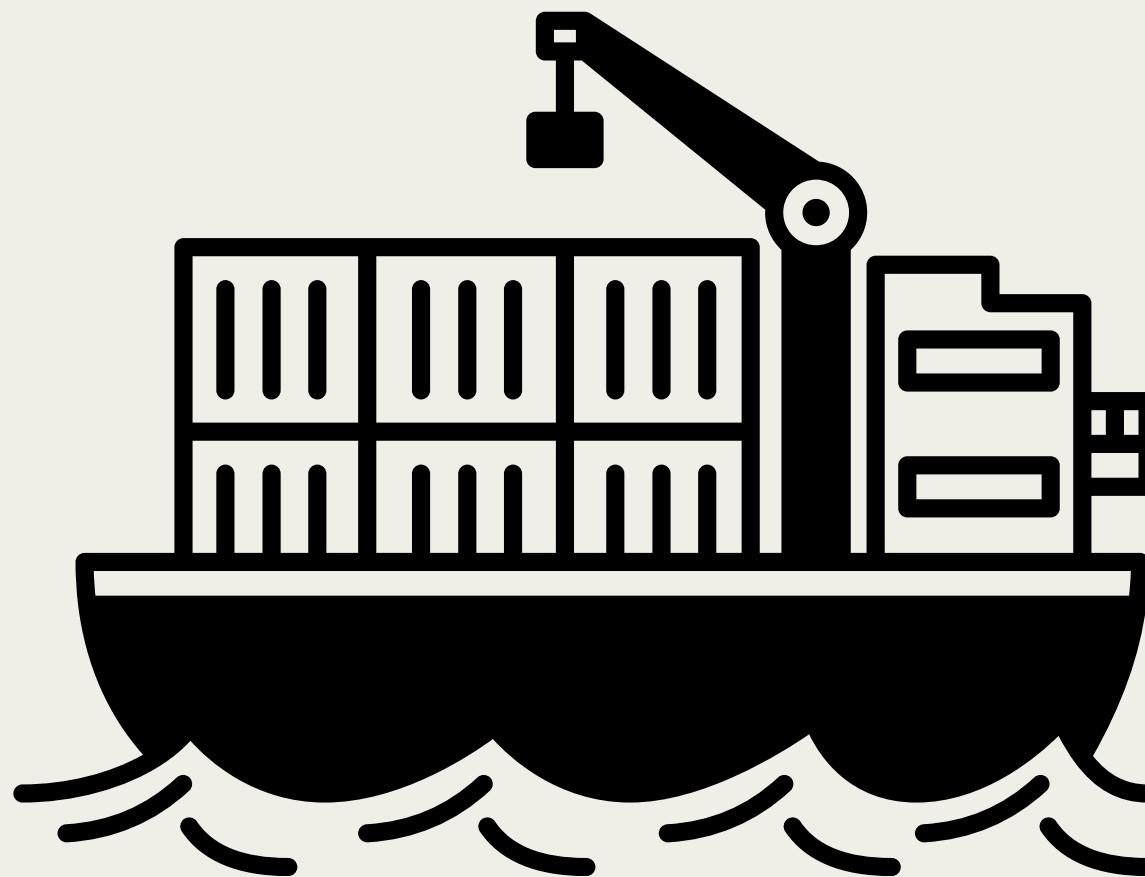
Which forms do I need to fill out?

What should be noted when transporting this category?

Are the current temperature conditions for the truck meeting the standards?

Is the current warehouse loading and unloading process correct for this category?

Non-specialized scenarios need Agent services



**Blue-collar minimum wage: \$16.50/hour
Annual income: \$68,640/year**

Blue-collar workers manually query data, which is time-consuming and labor-intensive.

Only large enterprises have the capability to develop specialized smart hardware for these scenarios.

Counterintuitive insight

Why are industries in the U.S. and China different, and why can't
SMBs produce them?

U.S. manufacturing is **slowly declining**

The barrier to entry for software development is decreasing

The barrier to hardware development

also needs to be lowered.

**Leveraging
developers to create
a competitive moat.**

Why build a developer ecosystem?

So developers don't have to reinvent the wheel themselves.

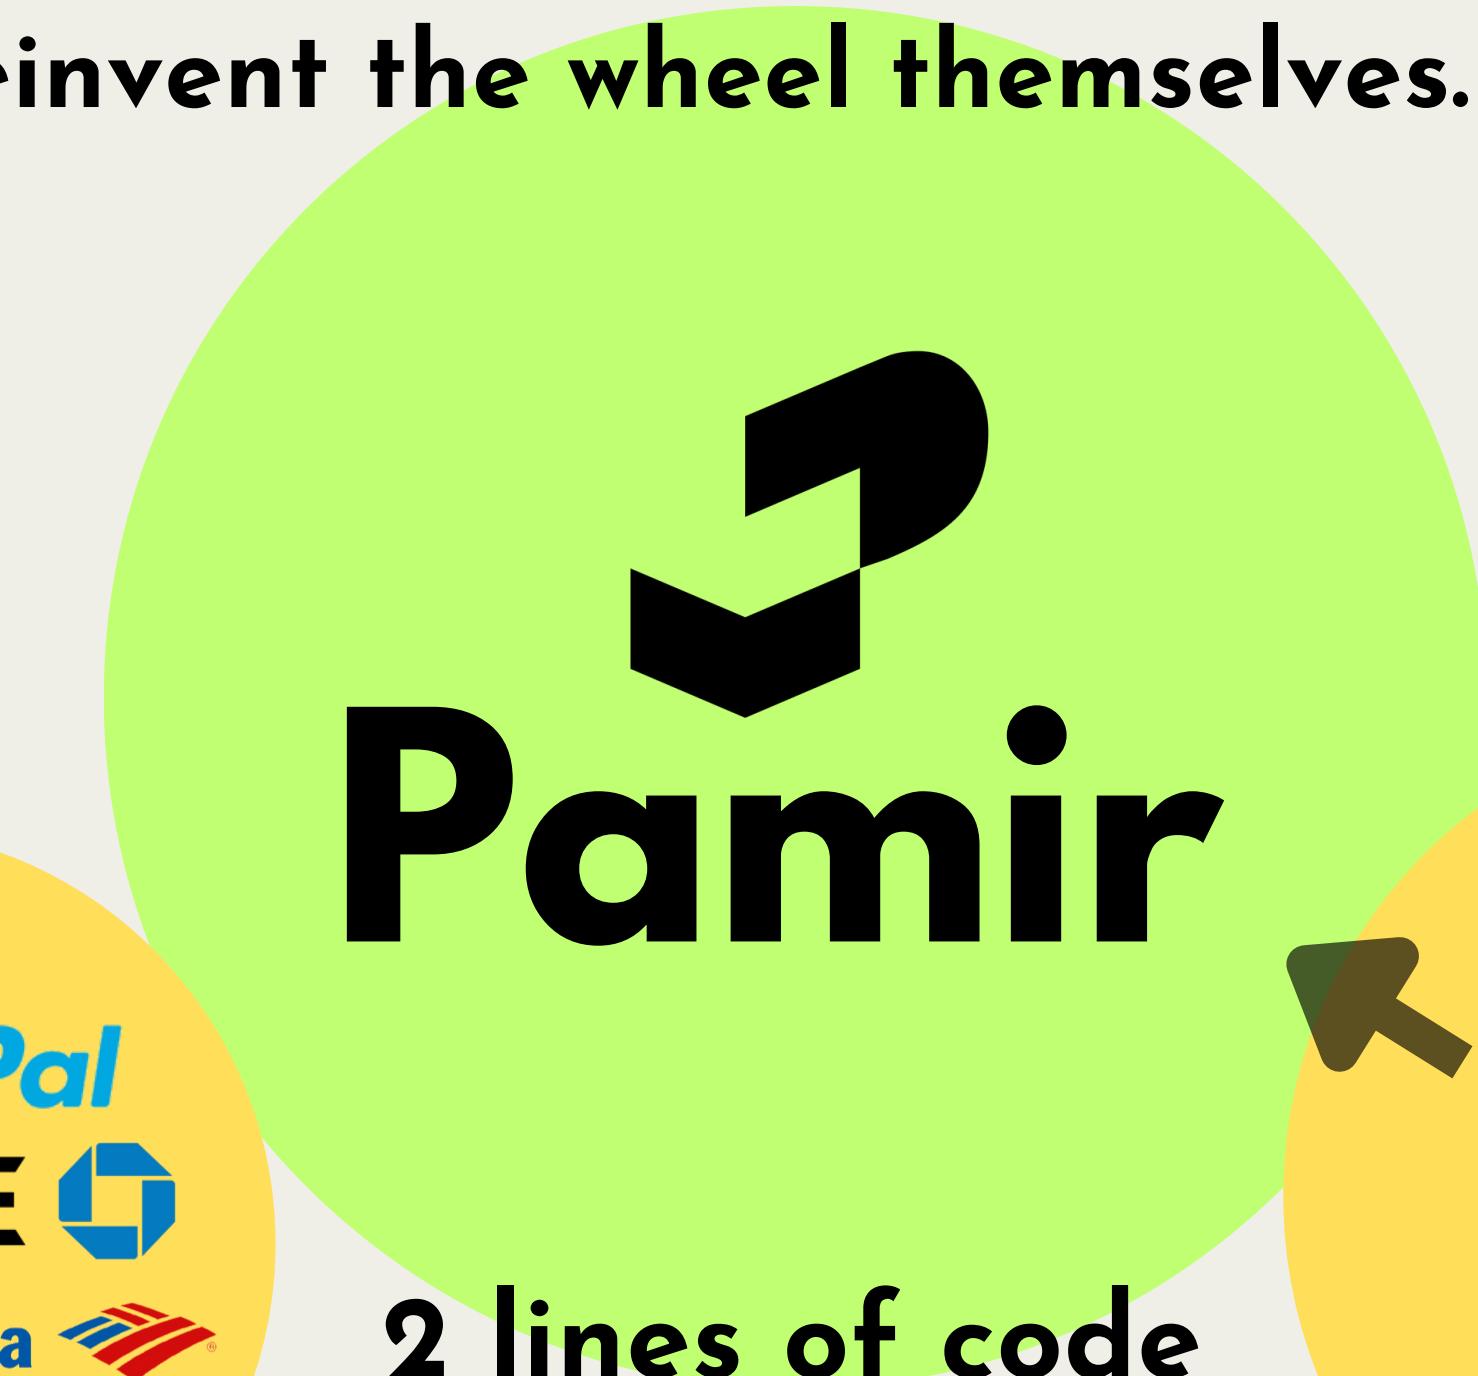


stripe

2 lines of code



Tens of thousands
of lines of code



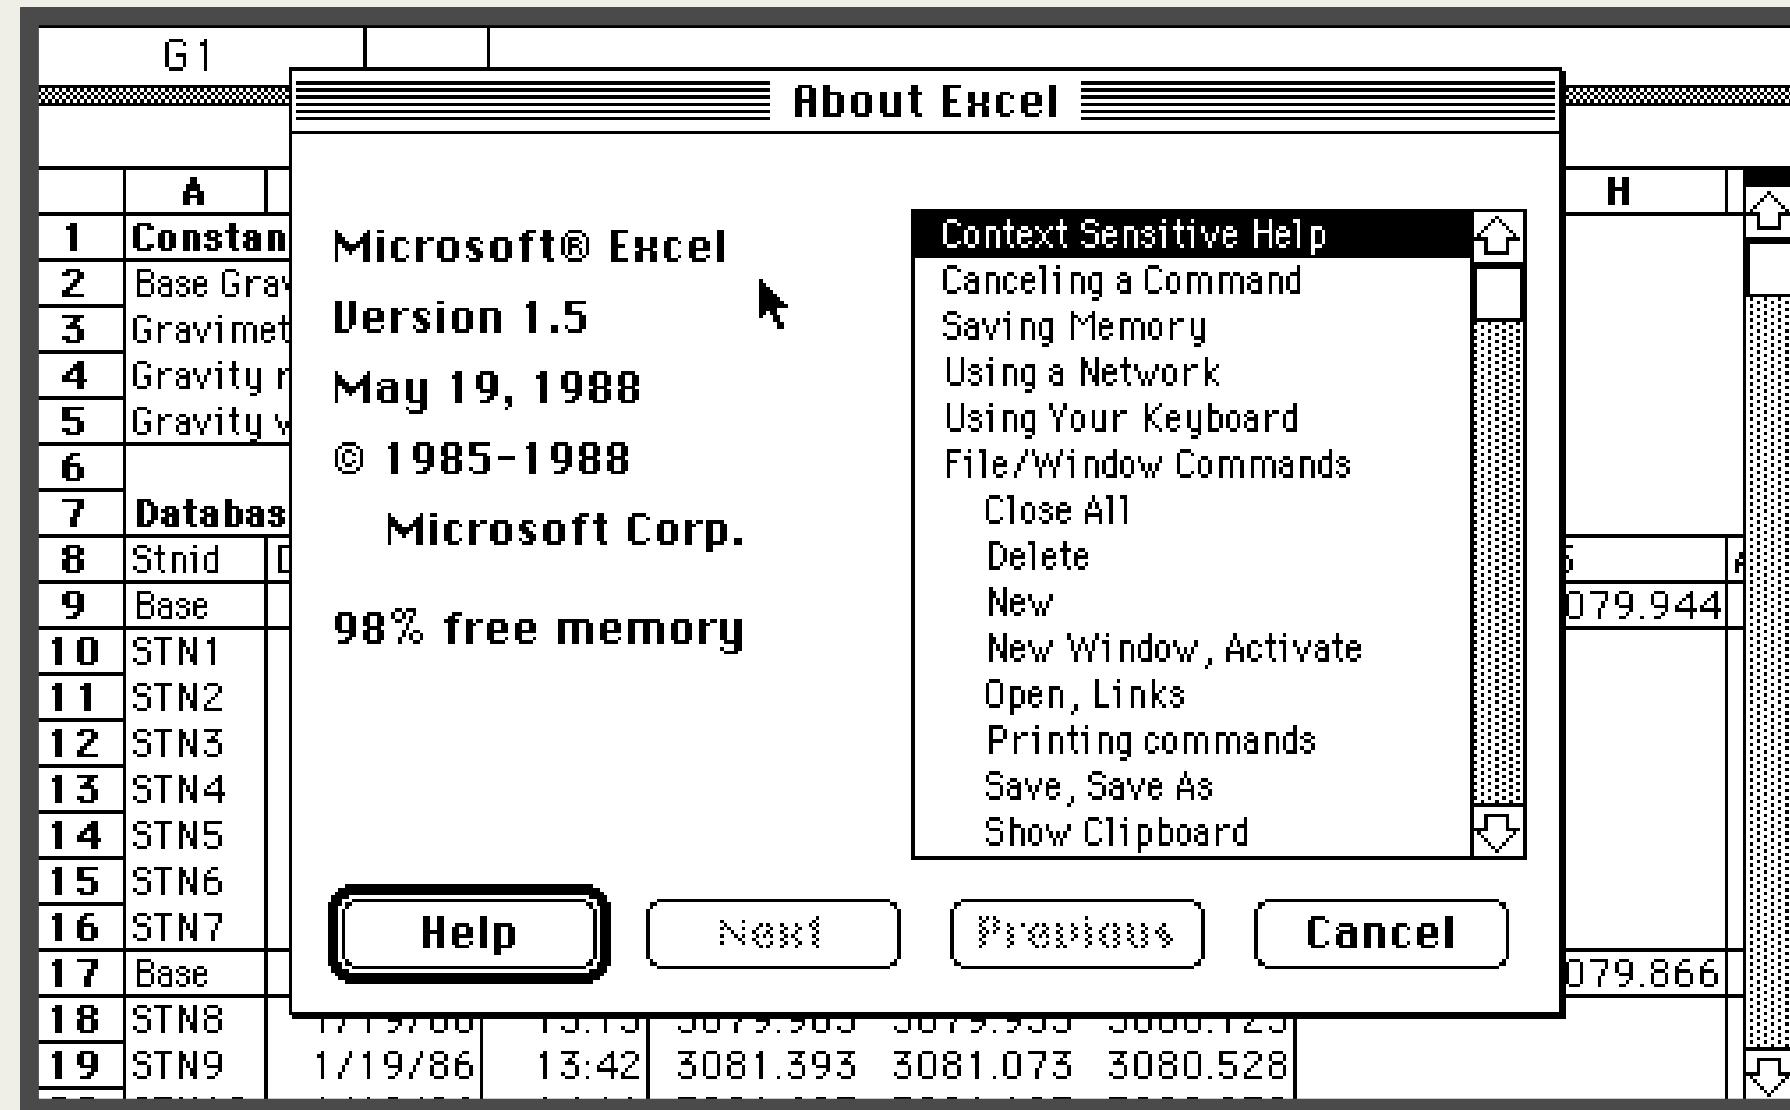
2 lines of code



Tens of thousands
of lines of code

**What was
implemented today?**

What happened on this day in the past? *spreadsheet on graphical interface*



*early days of excel :
allowed users to automate tasks through macros*

What should be done today? *spreadsheet on natural language interface*



allow users to automate decision through agentic data query

Business model

- \$250 / device
- \$50 / Month SaaS OTA service
- Set up a strategic department
to handle customized solutions
for large enterprises (toB).

Enterprise Purchase Logic

The monthly reduction in loss per device should be
≥ device cost / expected recovery cycle.

Taking logistic SMB as an example, conservatively, if the device can reduce the cargo loss rate from 1% to 0.7%, and assuming each task is worth 50 USD, with a device lifecycle of 3 years, each worker processes 2,500+ tasks annually, and there are 100 blue-collar workers:

The device could reduce cargo loss by approximately 37,500 USD annually.

Business Model: Different categories



**IQ 50: Focus on
monitoring systems**

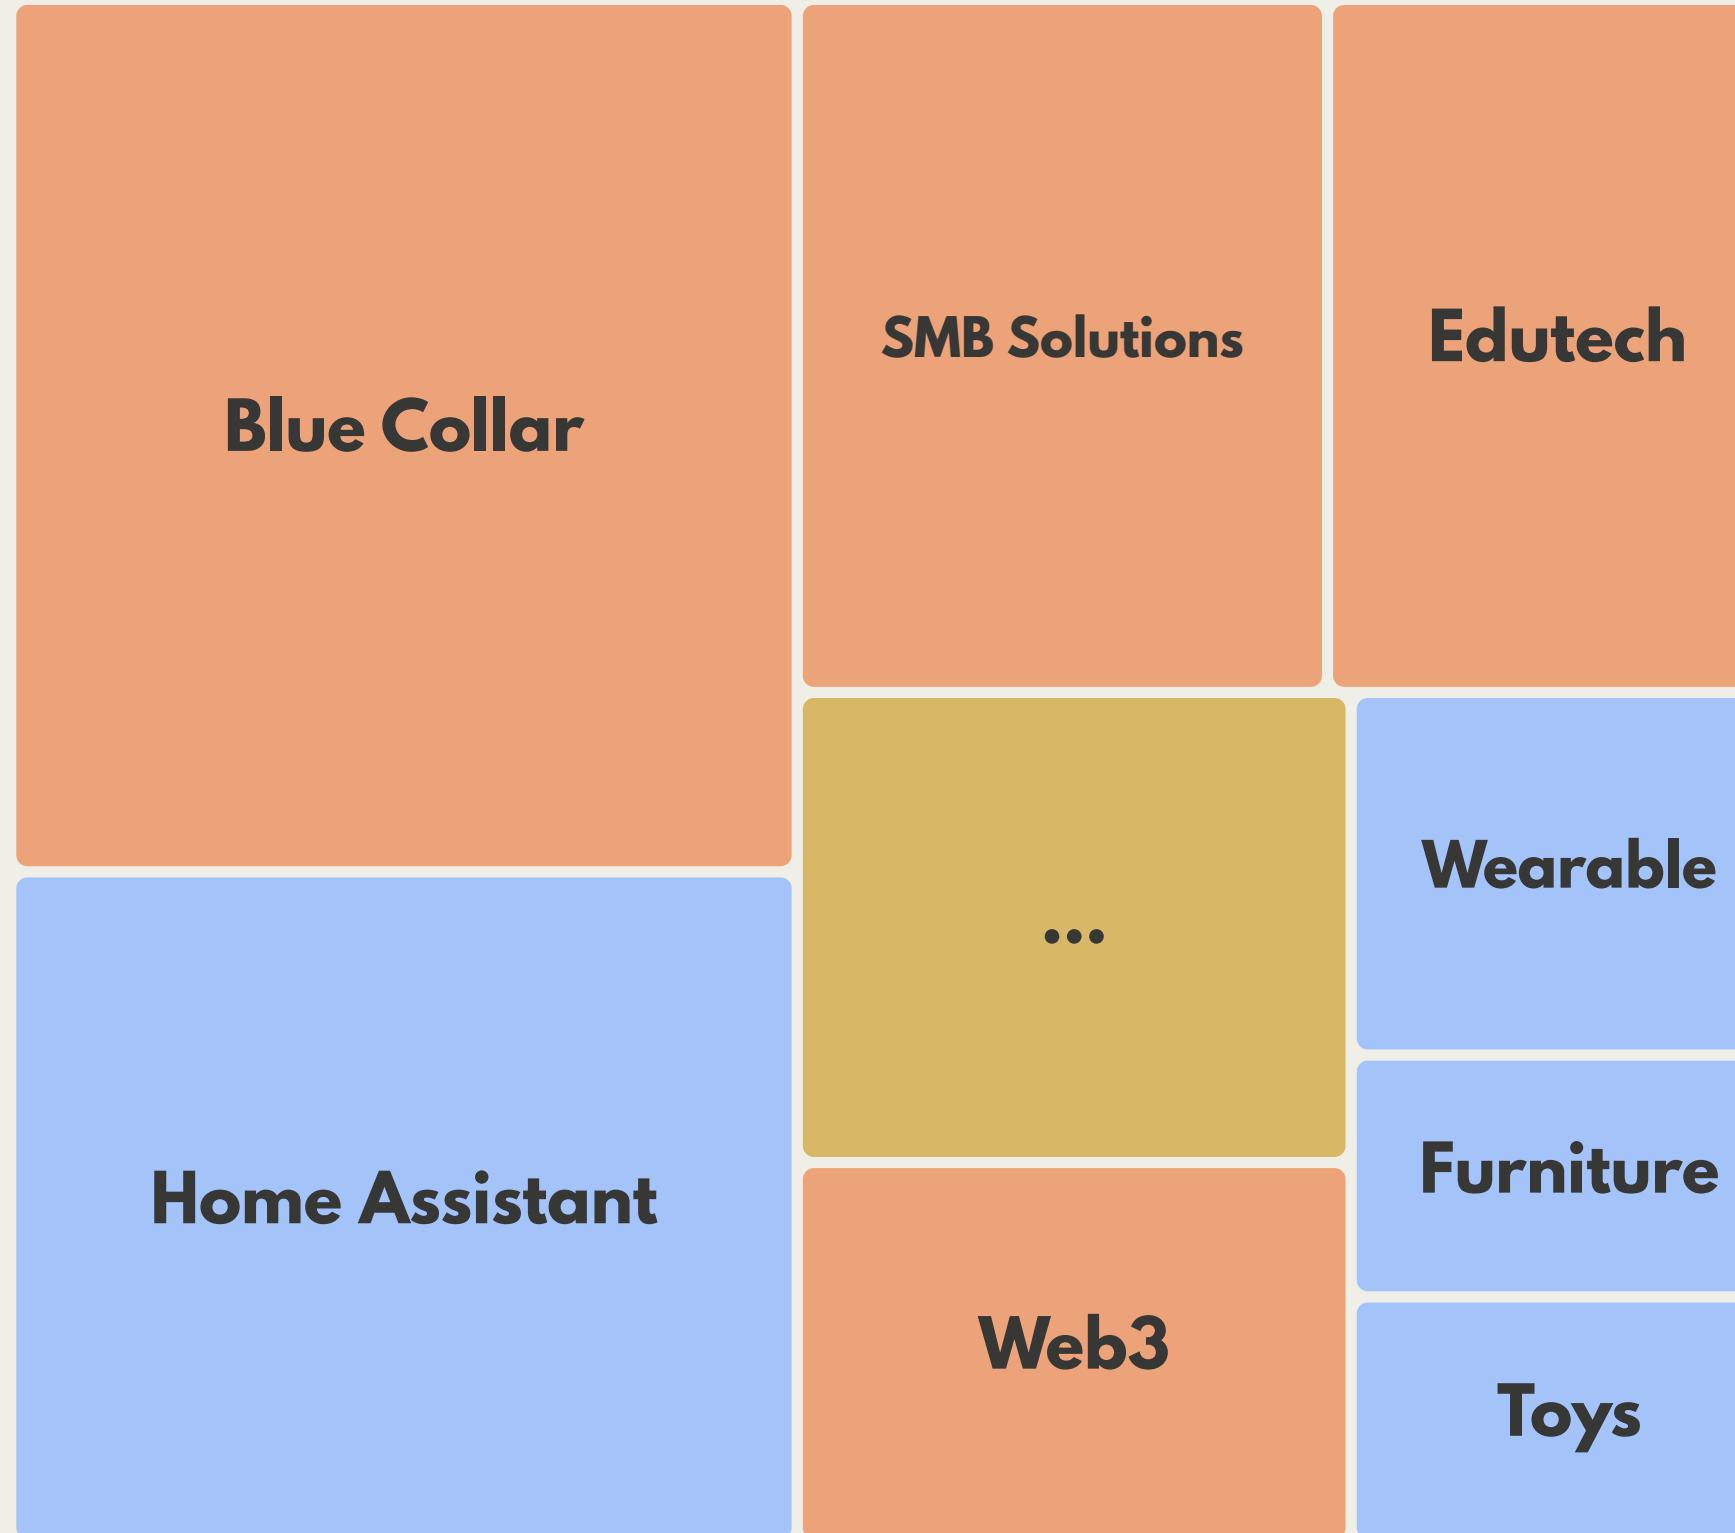


**IQ 100: Focus on
enterprise queries**



**IQ 150: Focus on management
decision-making**

First-mover advantage, but must also survive among competitors



ToC: Obvious market segments will be dominated by Shenzhen consumer hardware companies, leading to a race to the bottom.

ToB: In clear market segments, large companies will create independent vertical solutions, with the approach of "everyone just bring your own solution."

Competitors

**Model providers and chip manufacturers are not
our competitors.**

**Model providers currently lack the bandwidth to
focus on these tasks; they are more interested in
directly purchasing our solutions.**

Barriers / Advantages

Developer ecosystem

Standardized SDK interface

Continuous SaaS services

**The most cost-effective embedded
intelligence**

Vision

**Use IQ 100 for all embedded QA scenarios in the era
of natural language interaction.**

**Build an embedded Agent development platform for
the natural language interaction era.**

**The smallest box with the greatest intelligence.
Become the infrastructure of the future.**

TEAM

8-Year Classmates: Co-founded the DJI RoboMaster Club at Purdue University.



Qualcomm

- **Qualcomm Research Center** - Developed edge-side neural networks for autonomous driving, with 2 patents under the Qualcomm Innovation Fellowship.
- Qualcomm Autonomous Driving SoC Platform: Core team member for B2B business.
- Amazon SDE

 **Microsoft**

- **Microsoft Surface** Core R&D for compute devices.
- Extensive experience in consumer electronics development (0-1)
- Iterated through 5 generations of Surface product development

Appendix

Why Now

Competition Window

- NVIDIA/Qualcomm pushing efficient chips
- 158,000+ models on Hugging Face
- Yet no one focusing on Embedded Edge Agent for legacy industry.

Better Small Model

7B models match GPT-3.5 quality at 1/25th size

Better AI Chipset

Edge AI chips at \$99/12 TOPS sweet spot

Market Demand

140+ inbound calls in 30 days, 1M+ Views

Good Model

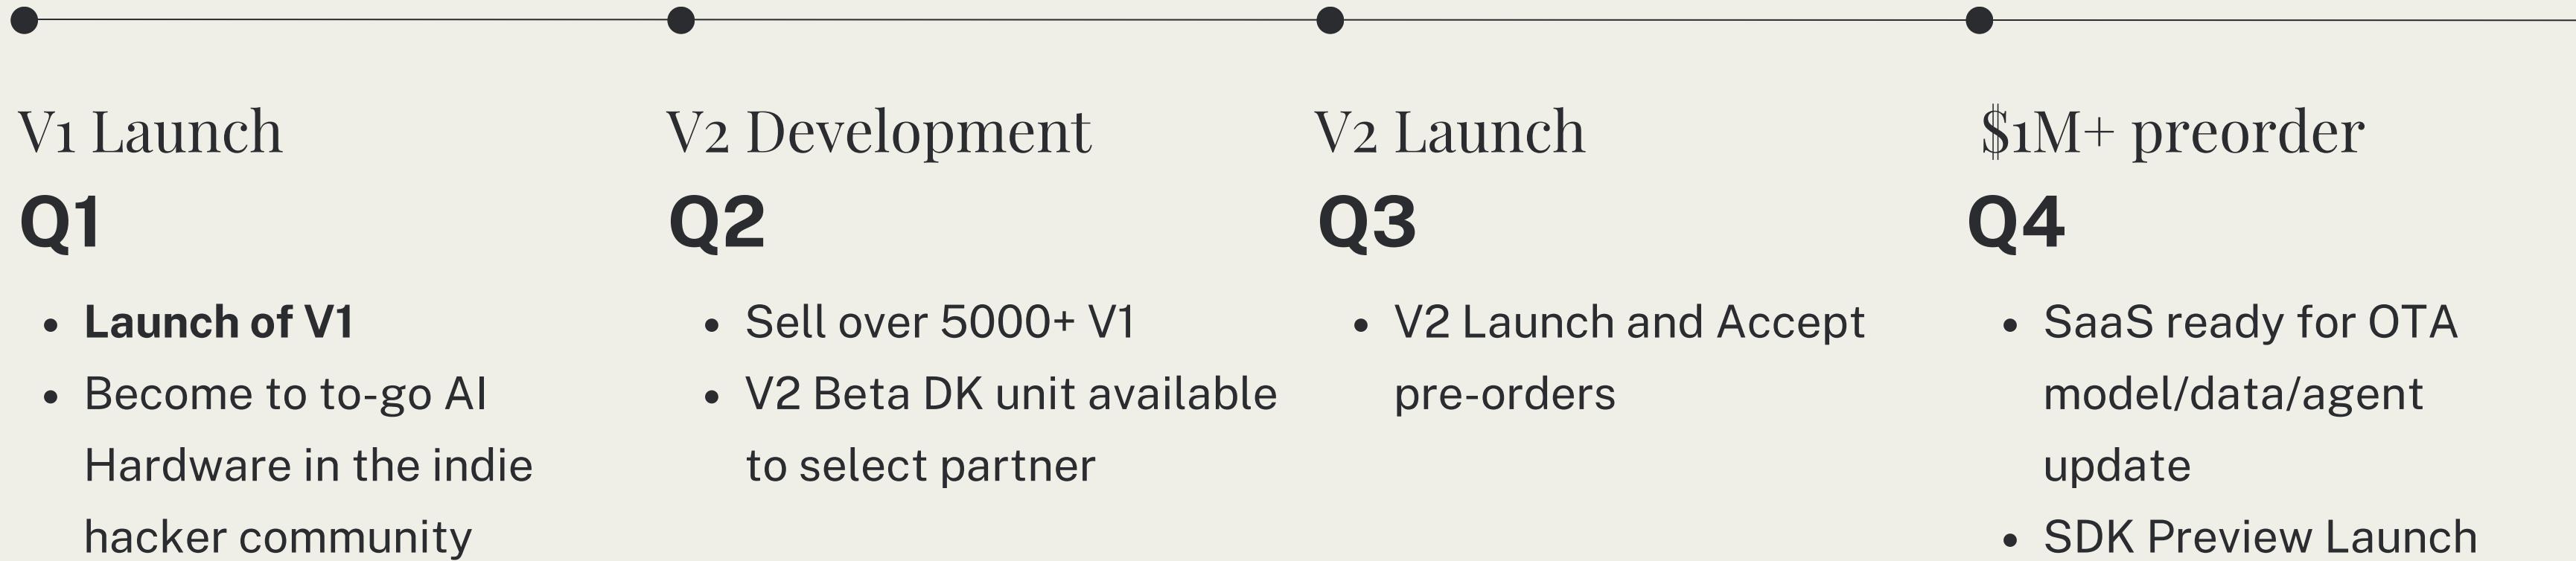
+

Efficient Hardware

+

Good RAG

Product roadmap



Real world



Sensor

embedded/SDK

display/audio/cam/...

interface

LLM SDK

Voice SDK

prompt engineering (for developer/consumer)

output modal

output SDK

external systems

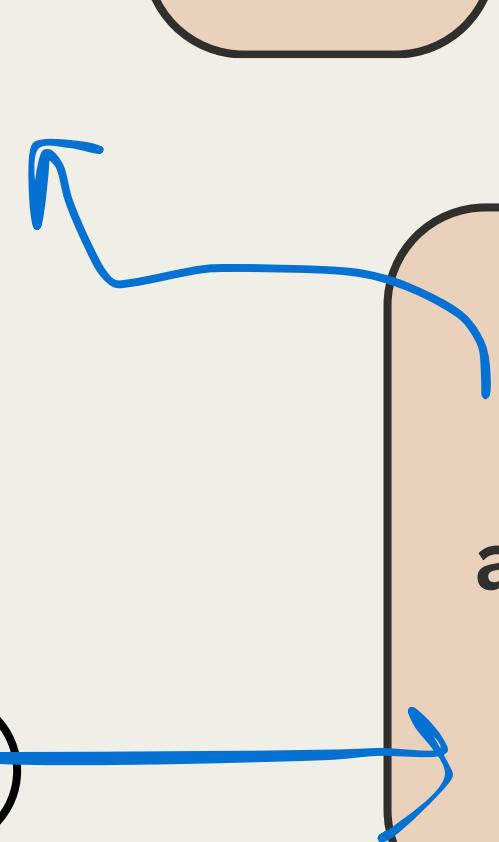
pro-sumers
new sensor or
new function
calls

mem/logs

RAG SDK

DB

agent framework



and how much we gonna make

2025Q4

10,000unit * \$250 = \$2.5M

10% SaaS conversion = \$120k ARR

2026Q2

50,000unit * \$250 = \$12.5M

30% SaaS conversion = \$1.8M ARR

GTM (to young developers)

Early Entry Points
Developers

Traditional enterprise data queries
Agent decision-making in blue-collar-intensive industries

GTM via Instagram, YouTube, TikTok (validated)
content creation via developer community
developer + branding = winning
check [flipperzero](#) for GTM reference

Progress (rockchip)

RAG to embed 15 sec for per 100 row csv data

LLM system : avg ~10 sec TTFT

- **soon next :**
 - **7B mix quant > 13token/s, TTFT 1.5s**
 - **3B you can expect > 20token/s**

Voice system : avg ~5sec inferencing time, outperform alexa level

Vision system : avg ~10sec response time

Power : 5-11 watts

Progress (Qualcomm)

2025 Jan Qualcomm 10 Dev Kit prototype

2025 March Qualcomm 200 early batch

2025 June DEFCON batch

2025 Q3 15+ SMB vertical integration

Long existing Demand:

Enterprise data + labor-intensive scenarios = better decision-making

Hi, we are members of the innovation office at [REDACTED], an engineering, construction, professional services and technology company. We are interested in using your devices for a hackathon with innovation-minded employees to encourage innovation in our employee-owned company. We have some potential use cases that we'd like to try especially in the field where connectivity is sometimes an issue.

We are [REDACTED], a company focused on developing cutting-edge Smart Cart-based retail solutions. In addition, we are engaged in advancing AI-driven robotics. We are seeking a partnership to further strengthen and enhance the robustness of our solutions. Your expertise in AI edge computing aligns perfectly with our vision, and we see significant potential for a successful collaboration.

Insight (product)

Edge-side Intelligence ≠ Mobile and PC Products

Users no longer expect to interact through screens, text, or by building their own apps!

Technical barriers will be lower, the complexity of technology will also reduce.

It's unrealistic to expect users to write scripts or bring their own solutions.

Hardware must be general-purpose intelligent.

Agents must be plug-and-play hardware that adapts firmware autonomously.

The intelligent agent should dynamically modify firmware, databases, and self-generate action code in real time, based on user natural language demands.

Insight (demand)

Most AI hardware companies are still **tunnel-visioned** on
ToC applications.

Blue-collar industries are where opportunities concentrate:
Characteristics: **Simple tasks, long working hours, high human error rates, requiring intelligence but not overly sophisticated.**

Opportunity: Simple QA scenarios are ideal for deploying early-stage intelligent agents but have been **overlooked**.

100 Years From Now

Interstellar cargo transportation will require extensive embedded intelligence systems.



Scavengers Reign: The interstellar transport ship DEMETER.